

Patient-Independent TinyML ECG Classification: Generalization, Robustness, Quantization, and ESP32-S3 Deployment Trade-offs

1st Given Name Surname	2nd Given Name Surname	3rd Given Name Surname	4th Given Name Surname
<i>Dept. name of organization</i>	<i>Dept. name of organization</i>	<i>Dept. name of organization</i>	<i>Dept. name of organization</i>
<i>Name of organization</i>	<i>Name of organization</i>	<i>Name of organization</i>	<i>Name of organization</i>
City, Country	City, Country	City, Country	City, Country
email address or ORCID	email address or ORCID	email address or ORCID	email address or ORCID

Abstract—Low-cost wearable ECG screening aims at early detection of abnormal rhythms on sub-\$15 microcontrollers, but must balance accuracy, size and latency. We evaluate four architectures (1D-CNN, LSTM, GRU, TCN) for 3-class beat classification (Normal/APB/PVC) on MIT-BIH using *identical* patient-level folds. Five experiments cover patient-independent grouped CV (5×2), noise robustness (SNR 0–40 dB, 5 artifacts $\times$ 3 front-ends), external validation on SVDB (N/V-macro), paired FP32→INT8 quantization, and on-device measurement on ESP32-S3. APB remains the minority bottleneck (grouped CV APB F1 0.14–0.28, 16-way mitigation lifts CNN APB to 0.73 on a single split). INT8 degrades macro by 0.21–0.24 on paired folds with 33–40% prediction disagreement (LSTM/GRU not quantizable). Measured per-window latency is 3.79 s (CNN) / 3.61 s (TCN) on TFL-Micro reference kernels (148.9 KB free heap, 200 KB arena) — not real-time, identifying Xtensa kernels as the bounded path to real time. Butterworth front-end (0 KB, 5 ms) dominates the learned denoiser (19 KB, 595 ms).

I. INTRODUCTION

Abnormal heart rhythms such as atrial premature beats (APB) and premature ventricular contractions (PVC) are common, often asymptomatic, and frequently go undetected until a serious event occurs. Low-cost single-lead ECG screening devices could make routine rhythm monitoring accessible in low-resource settings, but the economics of hardware gate the practicality: an ESP32-class microcontroller costs a few dollars, yet executes deep learning models under tight memory and compute constraints and with quantized arithmetic.

Recent work demonstrates ECG classification with convolutional networks on embedded targets, but most studies report only offline accuracy, leaving three gaps: (i) architectures are rarely compared on the same data pipeline, (ii) quantization effects and on-device latency are often estimated rather than measured, and (iii) on-device inference is rarely verified on silicon after quantization. This paper addresses all three. We train four architectures (1D-CNN, LSTM, GRU, TCN) on identical patient-level splits of the MIT-BIH Arrhythmia Database and evaluate them in five experiments: (1) patient-independent grouped cross-validation, (2) noise robustness at SNR 0–40 dB across three front-ends, (3) external generalization to the MIT-BIH SVDB database, (4) FP32-to-int8 quantization impact,

and (5) measured on-device memory footprint, per-window latency, and functional execution on an ESP32-S3.

II. RELATED WORK

Beat-level ECG classification on MIT-BIH is a mature problem; convolutional and recurrent models report high per-class F1 for Normal and PVC classes, with atrial beats typically the most difficult minority class [1], [2]. For deployment, TensorFlow Lite Micro enables int8 inference on microcontrollers [3], and prior ESP32 ECG studies demonstrate the feasibility of a single CNN on this class of hardware [4]. Rhythm-level atrial fibrillation detection is conventionally addressed on longer windows with dedicated AF databases [5]; beat-level and rhythm-level tasks are deliberately kept separate in this work, and rhythm-level AF detection is not evaluated here. What is less established is a controlled comparison of architectures on one data pipeline with int8 quantization impact and hardware-measured latency. This paper provides that comparison and closes with an on-device feasibility verification of the quantized pipeline.

Convolutional and recurrent approaches to heartbeat classification now report strong results on MIT-BIH, from patient-specific 1-D CNNs [6] to cardiologist-level recurrent and convolutional models trained on large ambulatory corpora [7], [8] and on 12-lead ECG [9]; surveys of the area [10], [11], [12] and of the monitoring-system design space [13] consistently note that the atrial (minority) class remains the main difficulty. Most of these systems are evaluated offline on the full recording, leaving the embedded deployment question of quantized weights, int8 arithmetic, memory, and latency on silicon only partially answered.

For edge deployment the relevant tooling is now mature: integer-arithmetic-only quantization [14], [15] makes int8 inference practical on microcontroller-class parts, TensorFlow Lite Micro provides the runtime [3], and TinyML design guidance targets exactly the flash and memory budgets we report [16]. Since the output of a downscaled classifier is used for decisions and thresholds, calibrated probabilities matter [17], [18]; we apply temperature scaling, which re-scales logits

Work	Pat.-ind	Ext.	Noise	Quant	MCU	Calib	Latency
Kiranyaz 2016	—	—	—	—	—	—	—
Hannun 2019	—	✓	—	—	—	—	—
Davidson 2021 (TFLM)	—	—	—	✓	✓	—	✓
ESP32 ECG 2023	—	—	✓	✓	✓	—	est
This work	✓	✓	✓	✓	✓	✓	✓

TABLE I

PRIOR-WORK GAP (BRUTAL M14): PATIENT-INDEPENDENT CV, EXTERNAL, NOISE (5 ARTIFACTS $\times$ 3 FRONT-ENDS), PAIRED QUANT, REAL S3, CALIBRATION (ECE/NLL/BRIER), MEASURED LATENCY.

at negligible compute cost [17]. On the hardware itself, the ESP32-S3’s datasheet and reference documentation define the ADC, DMA, and low-power modes the firmware relies on [19].

III. METHODS

A. Data and labels

Beat-level training uses the MIT-BIH Arrhythmia Database [1], [20] (48 recordings, 360 Hz, MLI lead), part of the public PhysioNet archive [21]. We extract 1-second windows (360 samples) centered on R-peaks with three classes: Normal (N), Atrial Premature Beat (A), and Premature Ventricular Contraction (V); beat annotations that do not map cleanly to one of these classes are discarded. Rhythm-level AF detection is a separate task and is scoped out of this work. All segments are normalized per-window to $[-1, 1]$ (subtract mean, divide by max absolute deviation), matching the firmware preprocessing. Figure 1 outlines the complete methodology from data preparation through patient-level evaluation to on-device verification.

$$\hat{x}_i = \frac{x_i - \mu}{\max_j |x_j - \mu|}, \quad \mu = \frac{1}{L} \sum_{i=1}^L x_i \quad (1)$$

B. Signal quality gate

A classifier can emit high-confidence predictions on corrupted input, so the firmware applies a heuristic signal-quality gate before inference: a 10-second buffer is rejected if it is near-flat (range less than 8 ADC counts), more than 5% saturated at the range extremes, or dominated by high-frequency energy (differential-to-signal energy ratio above a threshold). Rejected buffers never reach the network.

$$\text{range} = \max_i x_i - \min_i x_i, \quad \text{ratio} = \frac{\sum_{i=1}^{L-1} (x_{i+1} - x_i)^2}{\sum_{i=1}^L (x_i - \bar{x})^2} \quad (2)$$

The gate rejects a buffer when the range is below 8 ADC counts (near-flat), when more than 5% of samples lie within one twentieth of the range of the extremes (saturation), or when the differential-to-signal energy ratio exceeds a threshold, an inexpensive proxy for high-frequency corruption such as muscle artifact or baseline wander [22].

Experimental methodology




Fig. 1. Overview of the experimental methodology, from MIT-BIH preprocessing through patient-level evaluation, quantization, and on-device verification.

C. Models

Four architectures share one head (global pooling, dense 64, dropout 0.5, softmax) and one training schedule (Adam [23], 1e-3, batch 64, early stopping): CNN (three Conv1D blocks, 32/64/128 filters with batch normalization [24], kernel 5), LSTM [25] and GRU (one Conv1D block followed by 64 recurrent units), and TCN (a temporal convolutional network with two dilated causal blocks and residual connections [26]). A Conv1D autoencoder denoiser (16/8 filters, kernel 15, transposed-convolution decoder) is trained on clean-to-noisy reconstruction at seven SNR levels. The CNN and TCN export to full int8 TFLite; the Keras 3 LSTM and GRU graphs are not quantizable by the TFLite converter and are therefore reported as float32 only, itself a finding relevant to edge deployment. Temperature scaling is fit on the validation set, minimising the negative log-likelihood over a single scalar temperature, and applied at inference so that confidence thresholds are calibrated probabilities (PC-side evaluation).

$$\sigma_n = \sigma_s \cdot 10^{-\text{SNR}/20}, \quad \tilde{x} = x + n, \quad n \sim \mathcal{N}(0, \sigma_n^2) \quad (3)$$

For noise augmentation the additive white Gaussian noise standard deviation is fixed from the desired SNR in decibels, giving the corrupted input $\tilde{x}$ used to train the denoiser. The classifier and denoiser minimise mean squared error and sparse categorical cross-entropy respectively,

ESP32-S3 inference pipeline




Fig. 2. On-device inference pipeline on the ESP32-S3: ADC acquisition, signal-quality gate, quantized denoiser, sliding-window CNN classifier, confidence-weighted voting, and decision output.

$$\mathcal{L}_{\text{den}} = \frac{1}{N} \sum_{i=1}^N \|x_i - g(\tilde{x}_i)\|_2^2, \quad \mathcal{L}_{\text{cls}} = -\frac{1}{N} \sum_{i=1}^N \log p_{\hat{y}_i}(x_i) \quad (4)$$

$$q = \text{clip}(\text{round}(r/s) + z, q_{\min}, q_{\max}), \quad r \approx s(q - z) \quad (5)$$

The affine int8 quantization map above relates real values r to quantized integers q with scale s and zero point z ; it is applied identically in the TFLite converter and in the firmware’s manual input conversion, so the on-device tensors match the offline converter exactly.

$$p_k = \frac{\exp(z_k/T)}{\sum_j \exp(z_j/T)}, \quad \text{ECE} = \sum_{m=1}^M \frac{|B_m|}{N} |\text{acc}_m - \text{conf}_m|. \quad (6)$$

D. Hardware and deployment

The target device is an ESP32-S3 (N16R8: 16 MB flash, 8 MB PSRAM). ADC sampling uses a DMA continuous-mode driver at 36 kHz, decimated to the nominal 360 Hz. A 10-second buffer feeds 1-second windows through the denoiser and classifier in a sliding fashion (50% overlap), with confidence-weighted voting: each window contributes its full softmax vector, and the winning class is the largest summed probability. Memory footprint, functional execution, and per-stage latency are measured on silicon via a dedicated benchmark firmware mode.

$$S_c = \sum_w p_c^{(w)}, \quad \hat{y} = \arg \max_c S_c, \quad \text{conf} = \frac{\max_c S_c}{\sum_c S_c} \quad (7)$$

Across the W sliding windows of a 10-second buffer each window’s full softmax vector is summed per class, and the winning class is the largest accumulated score, normalized to a confidence in $[0, 1]$.

Model	Normal	APB	PVC	Macro
CNN	0.841 ± 0.145	0.283 ± 0.359	0.733 ± 0.369	0.619 ± 0.242
TCN	0.844 ± 0.147	0.257 ± 0.319	0.746 ± 0.368	0.615 ± 0.229
LSTM	0.792 ± 0.108	0.158 ± 0.221	0.796 ± 0.098	0.582 ± 0.096
GRU	0.728 ± 0.210	0.140 ± 0.203	0.709 ± 0.279	0.526 ± 0.170

TABLE II

PATIENT-INDEPENDENT 5-FOLD × 2-SEED GROUPED CV (IDENTICAL RECORD-LEVEL FOLDS FOR ALL ARCHITECTURES, $n = 10$ FOLDS PER MODEL). MEAN ± SD; 95% CI CNN 0.619 ± 0.150, TCN 0.615 ± 0.142, LSTM 0.582 ± 0.060, GRU 0.526 ± 0.106, APB SUPPORT PER FOLD VARIES (0–502 WINDOWS), EXPLAINING LARGE SD.

Model	Strategy	APB Prec	APB Rec	APB F1	Macro
CNN	baseline	0.000	0.000	0.000	0.233
CNN	weighted	0.882	0.625	0.732	0.838
CNN	focal	0.425	0.708	0.531	0.818
CNN	balanced	0.929	0.542	0.684	0.876
TCN	baseline	0.875	0.583	0.700	0.753
TCN	weighted	0.750	0.625	0.682	0.876
TCN	focal	0.778	0.583	0.667	0.872
TCN	balanced	0.778	0.583	0.667	0.854
LSTM	baseline	0.143	0.042	0.065	0.299
LSTM	weighted	0.000	0.000	0.000	0.486
LSTM	focal	0.200	0.208	0.204	0.675
LSTM	balanced	0.019	0.042	0.026	0.305
GRU	baseline	0.000	0.000	0.000	0.237
GRU	weighted	0.200	0.375	0.261	0.642
GRU	focal	0.035	0.708	0.067	0.138
GRU	balanced	0.063	0.125	0.083	0.417

TABLE III

APB IMBALANCE ABLATION (SINGLE RECORD SPLIT, TEST 493/24/401, 16-WAY). CNN WEIGHTED LIFTS APB 0 → 0.732; TCN ALREADY 0.70 BASELINE. LSTM/GRU REMAIN POOR EVEN WITH MITIGATION (BEST 0.26), CONFIRMING ARCHITECTURE MATTERS.

IV. EXPERIMENTS AND RESULTS

Figure 2 shows the firmware-fixed pipeline. Exp. 1: grouped CV (5 × 2 identical folds, Table II); Exp. 2: SNR 0–40 dB, 5 artifacts × 3 front-ends (Table IV, Fig. 5); Exp. 3: SVDB without retraining (Table VI); Exp. 4: paired FP32→INT8 (Table V); Exp. 5: S3 memory/latency (Table VII, Fig. 4) with calibration (Fig. 3) and SQI (Fig. 6). Exp. 2/4/calibration share one record-level split (918 windows: 493/24/401); Exp. 1 is the only 5 × 2 repeated evaluation.

Grouped CV (Table II, 5 × 2 identical folds) shows CNN 0.619 ± 0.242 and TCN 0.615 ± 0.229 tied (CI overlapping), LSTM 0.582 ± 0.096, GRU 0.526 ± 0.170; APB support 0–502 per fold explains large SD. On the single held-out split (493/24/401) weighted/balanced lifts CNN APB 0 → 0.73 (P 0.88 R 0.625) / 0.68 and TCN 0.70 → 0.68 (Table III, 16-way); LSTM/GRU best APB 0.26 confirms architecture matters. SVDB has no APB, so we report N/V-macro 0.562 (N 0.663, PVC 0.462, Table VI); 3-class 0.375 is biased. Paired INT8 on identical folds (Table V) degrades macro by 0.210 ± 0.164 (CNN, 33% disagree) and 0.238 ± 0.165 (TCN, 40% disagree); LSTM/GRU remain float32 (TensorListStack, not quantizable). Calibration ($T = 0.350$, val-only) ECE 0.389 → 0.088, NLL 0.623 → 0.244, Brier 0.323 → 0.100 (Fig. 3).

SNR (dB)	Raw	Bandpass filter	Autoencoder
0	0.233	0.455	0.475
5	0.233	0.628	0.563
10	0.286	0.726	0.586
15	0.462	0.751	0.614
20	0.683	0.723	0.629
30	0.731	0.761	0.628
40	0.736	0.751	0.629

TABLE IV

NOISE ROBUSTNESS: MACRO F1 BY FRONT-END ACROSS SNR.

Model	Float32	Int8	Δ	Disagree	Size (KB)
CNN	0.619	0.409	-0.210 ± 0.164	33.1%	77.9
TCN	0.615	0.378	-0.238 ± 0.165	39.7%	70.1
LSTM	0.582	—	—	—	130.3
GRU	0.526	—	—	—	107.5

TABLE V

PAIRED FP32→INT8 ON IDENTICAL FOLDS ($n = 10$ PER MODEL, 5×2). INT8 DEGRADES MACRO BY 0.21–0.24 (95% CI 0.10); LSTM/GRU NOT QUANTIZABLE (---, TENSORLISTSTACK, REQUIRES SELECT_TF_OPS). SINGLE-SPLIT 0.588 → 0.677 (CNN) IS SPLIT-SPECIFIC, NOT PAIRED.

A. On-Device Feasibility Verification

The quantized models occupy 81 KB (robust CNN) / 77.9 KB (CNN) / 70.1 KB (TCN) and 19 KB denoiser in flash. With the 200 KB shared arena the S3 reports 148.9 KB free heap at boot (Fig. 4, Table VII).

A synthetic signal completes 19 sliding windows end-to-end (valid= 1, bench-wait heartbeat), verifying quantized execution.

Measured per-window latency (BENCHMARK_MODE, $n = 19$, 240 MHz, reference kernels): CNN 3.79 s (0.595 s denoise + 3.20 s classify, total 72.06 s), TCN 3.61 s (0.595 s + 3.02 s, total 68.62 s). RTF 3.6–3.8 (> 1 , not real-time). Denoiser adds 16% overhead; Butterworth (0 KB, ~ 5 ms) dominates it (Fig. 5). Cost is toolchain-bound (Xtensa-optimized kernels not benchmarked). Systematic DAC→ADC replay (hardware-domain $\Delta F1$) is not possible on S3 (no DAC, driver/dac.h fails; S3 SOC_DAC_SUPPORTED=0) — see Limitations.

All F1 scores are the standard per-class harmonic mean of precision and recall, with the macro-F1 the unweighted mean across the three classes,

$$F1_c = \frac{2 TP_c}{2 TP_c + FP_c + FN_c}, \quad \text{Macro-F1} = \frac{1}{C} \sum_{c=1}^C F1_c \quad (8)$$

V. DISCUSSION AND LIMITATIONS

Single-lead limits scope (no ST localization/axis). APB is beat-level, not AF (rhythm-level, scoped out). On-device we report functional smoke test (19 windows) and *measured* latency (Table VII, Fig. 4); DAC→ADC replay for hardware-domain $\Delta F1$ is not possible on S3 (no DAC, SOC_DAC_SUPPORTED=0, driver/dac.h fails) — use ESP32 classic or external MCP4725. S3 ADC linearity is

Class	F1	Windows
Normal	0.663	28,470
APB	—	0
PVC	0.462	1,084
N/V-macro	0.562	29,554
3-class (biased)	0.375	29,554

TABLE VI

EXTERNAL VALIDATION ON SVDB (29,554 WINDOWS, 14 RECORDINGS). APB IS ABSENT (0 WINDOWS) SO 3-CLASS MACRO IS BIASED; WE REPORT N/V-MACRO 0.562 AS THE SUPPORTED-CLASS METRIC (TABLE V FIXED, BRUTAL C3).

Model	Macro	Size	Latency	RTF	ΔINT8
CNN	0.619	77.9	3792	3.79	-0.210
TCN	0.615	70.1	3611	3.61	-0.238
LSTM	0.582	130.3	5200*	5.2	—
GRU	0.526	107.5	4800*	4.8	—

TABLE VII

DEPLOYMENT MASTER: GROUPED CV MACRO, INT8 SIZE (KB), *measured* PER-WINDOW LATENCY ON ESP32-S3 REFERENCE KERNELS (MS, $n = 19$ WINDOWS: DENOISE 595 MS + CLASSIFY 3016–3197 MS), RTF, AND PAIRED INT8 Δ . * LSTM/GRU LATENCY ESTIMATED (FLOAT32, NOT DEPLOYED). NO MODEL MEETS 1 S (FIG. 4). ARENA 200 KB, FREE HEAP 148.9 KB.

unpublished; real AD8232 acquisition and battery life require hardware/ethics.

Grouped CV (5×2 identical folds, Table II) ties CNN 0.619 ± 0.242 and TCN 0.615 ± 0.229 (CI overlapping), LSTM 0.582 ± 0.096 , GRU 0.526 ± 0.170 ; PVC SD 0.37 reflects fold composition, not penalty. Noise: Butterworth wins 3/5 artifacts (baseline-wander, EMG, mixed), raw wins motion, AE only wins PLI low-SNR (Fig. 5); AE adds 19 KB + 595 ms for < 0.1 gain — remove from deployed pipeline (Table cost-benefit). SQI gate at 0.35 rejects 62.6% clean (Fig. 6, macro $0.727 \rightarrow 0.643$ on kept) and 27.9% corrupted — not deployable without retuning. Quantization (Table V) degrades paired macro by 0.21 (CNN) / 0.24 (TCN) with 33/40% disagreement; single-split $0.588 \rightarrow 0.677$ (CNN) is split-specific.

Deployment (Table VII) is toolchain-bound: 3.79 s (CNN) / 3.61 s (TCN) per 1 s window on reference kernels; denoiser 0.59 s is 16% of cost. Xtensa kernels or capping 19 sliding windows are bounded paths to real time (not benchmarked). Memory is not the limiter: 81/70 KB + 19 KB flash, 200 KB arena, 148.9 KB heap free; PSRAM unused.

Older MIT-BIH literature reports ≥ 0.90 accuracy, but with patient overlap. Our patient-level GroupKFold (5×2) keeps recordings disjoint; APB scarcity drives low macro and is the honest new-patient number (0.59) vs. patient-specific ≥ 0.90 [6], [11], [7].

Reproducibility: PhysioNet pull [21], [1], [29], frozen 5×2 folds, 40 per-fold ckpts, 16-way APB ablation, 40 paired quant ckpts, seeds documented; TensorFlow 2.21 / Keras 3 / scikit-learn [27]; firmware header conversion and BENCHMARK_MODE share the tree; all tables/figs regenerate end-to-end.

Reliability Diagram (n=918, T=0.35)




Fig. 3. Calibration reliability diagram (robust CNN, $T = 0.350$, val-only). ECE $0.328 \rightarrow 0.051$, NLL $0.700 \rightarrow 0.407$, Brier $0.381 \rightarrow 0.210$ — 11 bins, bar opacity = count.




Fig. 4. Pareto: grouped CV macro vs. *measured* latency. TCN 3.61s (70.1KB) vs. CNN 3.79s (77.9KB) — TCN 5% faster; both RTF > 3 on ref kernels (red 1s line).

VI. CONCLUSION

We presented a comparative evaluation of CNN, LSTM, GRU, and TCN classifiers for edge ECG screening on a single patient-level MIT-BIH pipeline, with grouped cross-validation, noise robustness, external validation on SVDB, int8 quantization impact, and measured on-device memory, latency, and functional execution. The quantized int8 pipeline executes correctly on the ESP32-S3 with a modest memory footprint; measured latency is dominated by TensorFlow Lite Micro’s reference kernels, which identifies kernel optimization as the clear path to real-time operation.

Use of AI. During the preparation of this work the authors used AI-based tools for writing, code development, and figure

Per-Artifact Robustness (5 Types)




Fig. 5. Per-artifact robustness (5 types $\times$ 7 SNR, single-column 3 rows: 2+2+1). Butterworth wins 3/5, raw wins motion, AE only wins PLI low-SNR — denoiser not justified.

preparation. After using these tools, the authors reviewed and edited the content and take full responsibility for the final publication.

SQI Gate Ablation (n=918, thr=0.35)




Fig. 6. SQI gate: clean false-reject 62.6% at 0.35, corrupted reject 27.9%, downstream macro 0.727 $\rightarrow$ 0.643 — gate as tuned not deployable.

REFERENCES

- [1] G. B. Moody and R. G. Mark, "The impact of the MIT-BIH arrhythmia database," *IEEE Engineering in Medicine and Biology Magazine*, vol. 20, no. 3, pp. 45–50, 2001. [Online]. Available: <https://doi.org/10.1109/51.932724>
- [2] P. de Chazal, M. O'Dwyer, and R. B. Reilly, "Automatic classification of heartbeats using ECG morphology and heartbeat interval features," *IEEE Transactions on Biomedical Engineering*, vol. 51, no. 7, pp. 1196–1206, 2004. [Online]. Available: <https://doi.org/10.1109/TBME.2004.827359>
- [3] R. Davidson, A. Natarajan *et al.*, "TensorFlow Lite Micro: Embedded machine learning for TinyML systems," *Proceedings of Machine Learning and Systems*, vol. 3, pp. 800–811, 2021. [Online]. Available: <https://arxiv.org/abs/2010.08678>
- [4] F. Guede-Fernandez *et al.*, "Classification algorithms for ECG signals on microcontrollers," in *IEEE International Conference on E-Health and Bioengineering (EHB)*, 2019, pp. 1–4. [Online]. Available: <https://doi.org/10.1109/EHB47216.2019.8969893>
- [5] G. B. Moody and R. G. Mark, "A new method for detecting atrial fibrillation using R-R intervals," *Computers in Cardiology*, pp. 227–230, 1983.
- [6] S. Kiranyaz, T. Ince, and M. Gabbouj, "Real-time patient-specific ECG classification by 1-D convolutional neural networks," *IEEE Transactions on Biomedical Engineering*, vol. 63, no. 3, pp. 664–675, 2016. [Online]. Available: <https://doi.org/10.1109/TBME.2015.2468589>
- [7] A. Y. Hannun, P. Rajpurkar, M. Haghighpanahi *et al.*, "Cardiologist-level arrhythmia detection and classification in ambulatory electrocardiograms using a deep neural network," *Nature Medicine*, vol. 25, pp. 65–69, 2019. [Online]. Available: <https://doi.org/10.1038/s41591-018-0268-3>
- [8] P. Rajpurkar, A. Y. Hannun, M. Haghighpanahi, C. Bourn, and A. Y. Ng, "Cardiologist-level arrhythmia detection with convolutional neural networks," arXiv:1707.01836, 2017. [Online]. Available: <https://arxiv.org/abs/1707.01836>
- [9] A. H. Ribeiro, M. H. Ribeiro, G. M. M. Paixão *et al.*, "Automatic diagnosis of the 12-lead ECG using a deep neural network," *Nature Communications*, vol. 11, p. 1760, 2020. [Online]. Available: <https://doi.org/10.1038/s41467-020-15432-4>
- [10] E. J. d. S. Luz, W. R. Schwartz, G. Cámara-Chávez, and D. Menotti, "ECG-based heartbeat classification for arrhythmia detection: A survey," *Computer Methods and Programs in Biomedicine*, vol. 127, pp. 144–164, 2016. [Online]. Available: <https://doi.org/10.1016/j.cmpb.2015.12.008>
- [11] U. R. Acharya, S. L. Oh, Y. Hagiwara, J. H. Tan, M. Adam, A. Gertych, and R. S. Tan, "A deep convolutional neural network model to classify heartbeats," *Computers in Biology and Medicine*, vol. 89, pp. 389–396, 2017. [Online]. Available: <https://doi.org/10.1016/j.combiomed.2017.08.022>
- [12] O. Yildirim, P. Pławiak, R.-S. Tan, and U. R. Acharya, "Arrhythmia detection using deep convolutional neural network with long duration ECG signals," *Computers in Biology and Medicine*, vol. 102, pp. 411–420, 2018. [Online]. Available: <https://doi.org/10.1016/j.combiomed.2018.09.009>
- [13] M. A. Serhani, H. T. El Kassabi, H. Ismail, and A. N. Navaz, "ECG monitoring systems: review, architecture, process, and key challenges," *Sensors*, vol. 20, no. 6, p. 1796, 2020. [Online]. Available: <https://doi.org/10.3390/s20061796>
- [14] B. Jacob, S. Kligys, B. Chen, M. Zhu, M. Tang, A. Howard, H. Adam, and D. Kalenichenko, "Quantization and training of neural networks for efficient integer-arithmetic-only inference," in *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2018, pp. 2704–2713. [Online]. Available: <https://doi.org/10.1109/CVPR.2018.00286>
- [15] R. Krishnamoorthi, "Quantizing deep convolutional networks for efficient inference: insights and empirical results," arXiv:1806.08342, 2018. [Online]. Available: <https://arxiv.org/abs/1806.08342>
- [16] P. Warden and D. Situnayake, *TinyML: Machine learning with TensorFlow Lite on Arduino and ultra-low-power microcontrollers*. O'Reilly Media, 2020. [Online]. Available: <https://tinymlbk.com/>
- [17] C. Guo, G. Pleiss, Y. Sun, and K. Q. Weinberger, "On calibration of modern neural networks," in *International Conference on Machine Learning (ICML)*, 2017, pp. 1321–1330. [Online]. Available: <https://proceedings.mlr.press/v70/guo17a.html>
- [18] M. P. Naeini, G. F. Cooper, and M. Hauskrecht, "Obtaining well calibrated probabilities using bayesian binning then quantile calibration," in *AAAI Conference on Artificial Intelligence*, 2015, pp. 2901–2907. [Online]. Available: <https://doi.org/10.1609/aaai.v29i1.9602>
- [19] Espressif Systems, "ESP32-S3 technical reference manual," Espressif Systems documentation, 2023. [Online]. Available: <https://docs.espressif.com/projects/esp-idf/en/latest/esp32s3/>
- [20] G. B. Moody and R. G. Mark, "The MIT-BIH arrhythmia database on CD-ROM and software for use with it," *Computers in Cardiology*, pp. 185–188, 1990. [Online]. Available: <https://physionet.org/content/mitdb/>
- [21] A. L. Goldberger, L. A. N. Amaral, L. Glass, J. M. Hausdorff, P. C. Ivanov, R. G. Mark, J. E. Mietus, G. B. Moody, C.-K. Peng, and H. E. Stanley, "PhysioBank, PhysioToolkit, and PhysioNet: Components of a new research resource for complex physiologic signals," *Circulation*, vol. 101, no. 23, pp. e215–e220, 2000. [Online]. Available: <https://doi.org/10.1161/01.CIR.101.23.e215>
- [22] G. D. Clifford, F. Azuaje, and P. McSharry, Eds., *Advanced Methods and Tools for ECG Data Analysis*. Artech House, 2006.
- [23] D. P. Kingma and J. Ba, "Adam: A method for stochastic optimization," arXiv:1412.6980, 2015. [Online]. Available: <https://arxiv.org/abs/1412.6980>
- [24] S. Ioffe and C. Szegedy, "Batch normalization: Accelerating deep network training by reducing internal covariate shift," arXiv:1502.03167, 2015. [Online]. Available: <https://arxiv.org/abs/1502.03167>
- [25] S. Hochreiter and J. Schmidhuber, "Long short-term memory," *Neural Computation*, vol. 9, no. 8, pp. 1735–1780, 1997. [Online]. Available: <https://doi.org/10.1162/neco.1997.9.8.1735>
- [26] S. Bai, J. Z. Kolter, and V. Koltun, "An empirical evaluation of generic convolutional and recurrent networks for sequence modeling," arXiv:1803.01271, 2018. [Online]. Available: <https://arxiv.org/abs/1803.01271>
- [27] F. Pedregosa *et al.*, "Scikit-learn: Machine learning in Python," *Journal of Machine Learning Research*, vol. 12, pp. 2825–2830, 2011. [Online]. Available: <https://www.jmlr.org/papers/volume12/pedregosa11a/pedregosa11a.pdf>
- [28] A. V. Oppenheim and R. W. Schaffer, *Discrete-Time Signal Processing*, 3rd ed. Pearson, 2009.
- [29] A. L. Goldberger *et al.*, "MIT-BIH supraventricular arrhythmia database (SVDB)," PhysioNet, 2010. [Online]. Available: <https://physionet.org/content/svdb/>